Lecture 3: Item frequency estimation in streams

Lecturer: Edo Liberty

Warning: *This note may contain typos and other inaccuracies which are usually discussed during class. Please do not cite this note as a reliable source. If you find mistakes, please inform me.*

Say we are given a stream of elements $X = [x_1, \ldots, x_N]$ where $x_i \in \{a_1, \ldots, a_n\}$. Let n_i denote the number of times element a_i appeared in the stream, i.e., $f_i = |\{j|x_j = a_i\}|$. Our goal is to estimate f_i for all frequent elements. This can be solved exactly by keeping a counter for each element $\{a_1, \ldots, a_n\}$. Alas, this might require, $\Theta(n)$ memory. Here we look for methods to approximate the values on f_i using $o(n)$ memory.

1 Sampling

The first and simplest approach is to use the uniform sampling approach above. That is, the algorithm draws samples uniformly at random from the stream with probability ℓ/N. Using Chernoff along with the union bound indicates that $\ell \in O(\log(n/\delta)/\varepsilon^2)$ is sufficient. Applying the union bound more carefully reduces the failure probability and therefore reduces ℓ, the expected number of samples.

2 Count Min-Sketches

Note that the space dependence of random sampling on ε is inversely quadratic which might be problematic for small values of ε. Count-Min sketches were introduced in [1][2] in two similar variants. They reduce the space complexity dependence on ε to only $1/\varepsilon$. The creation of the sketch is given in Algorithm **??**. The notation is that $h_1, \ldots, h_t$ are hash functions from the space of elements to the integers $[\lceil 2/\varepsilon \rceil]$.

Algorithm 1 Count Min Sketch: Add

    Input: ε, A
    $t \leftarrow \lceil \log(n/\delta) \rceil$, $b \leftarrow \lceil 2/\varepsilon \rceil$
    $C \leftarrow$ all zeros matrix of size $t \times b$
    for $i \in [N]$ **do**
      for $j \in [t]$ **do**
        $C[j, h_j(A_i)] = C[j, h_j(A_i)] + 1$
      end for
    end for
    Return: C

Algorithm 2 Count Min Sketch: Query

    Input: C, a
    Return: $\min_{j=1,\ldots,t} C[j, h_j(a)]$

To see why this works consider only one row of the sketch matrix. The value of $C[1, a]$ contains the frequency of a but also the sum of frequencies of all other items b for which $h_1(b) = h_1(a)$. Since the event that $h_1(b) = h_1(a)$ happens with probability $\varepsilon/2$ we have $\mathbb{E}[C[1, a] - f_a] \leq N\varepsilon/2$ by linearity of

expectation. By Markov's inequality we have that $\Pr[C[1,a] - f_a \geq \varepsilon N] \leq 1/2$. Therefore, any row in C provides a good approximate count for a with probability at least $1/2$. Since we return the minimal value of $\log[n/\delta]$ such estimates our failure probability reduces to δ/n. Using the union bound we get that all items receive a good approximation with probability at least $1 - \delta$. Note that we get the same guaranties as in the sampling solution but the space requirement reduced from $O(\log(n/\delta)\varepsilon^2)$ to $O(\log(n/\delta)/\varepsilon)$. Alas, the update time increases from $O(1)$ to $O(\log(n/\delta))$. In the next section we see how this can be improved and even derandomized.

3 Frequent Items

The item frequency approximation problem a brilliantly simple and deterministic algorithm in [3]. This algorithm was later rediscovered independently by both [4] and [5] who also improved its update time complexity. Their algorithm reduces the space requirement from $O(\log(n/\delta)/\varepsilon)$ to $O(1/\varepsilon)$. Their algorithm is given in Algorithm box 3. To prove the algorithm's correctness, let n' denote the sum of all counters

---

Algorithm 3 Lossy counting

> **Input:** $\varepsilon \in (0,1]$, A
> $\ell \leftarrow \lceil 1/\varepsilon \rceil$
> $C \leftarrow$ empty map from a to the integers with returned default value 0
> **for** $i \in [N]$ **do**
> $C[A_i] = C[A_i] + 1$
>   **if** $size(C) = \ell$ **then**
>   **for** $a \in C$ **do**
> $C[a] = C[a] - 1$
>   **if** $C[a] = 0$ **then**
> $del(C[a])$
>   **end if**
>   **end for**
>   **end if**
> **end for**
> **Return:** C

---

in the returned map C. Let $\delta_i = 1$ if the inner loop of the algorithm is executed in the i'th iteration and zero else. Note that in each iteration the sum of counters is increased by 1 and reduced by $\ell\delta_i$. Therefore $N' = \sum_{i=1}^{N} 1 - \ell\delta_i = N - \ell\sum_{i=1}^{N}\delta_i$. This gives that $\sum_{i=1}^{N}\delta_i \leq (N - N')/\ell \leq \varepsilon(N - N')$. Since $N' \geq 0$ and any single item counter is decreased at most $\sum_{i=1}^{N}\delta_i$ times we get that $f_a \geq C[a] \geq f_a - \varepsilon N$.

This reduces the amount of memory from $O(\log(n/\delta)/\varepsilon)$ required by Count-Min sketches to $O(1/\varepsilon)$. Moreover, some modifications to the data structure in the algorithm [5] allow updates to require only $O(1)$ operations. This significantly improves on the $O(\log(n/\delta))$ operations required by Count-Min sketches. As a last remark, note that this algorithm is deterministic which eliminates the failure probability altogether.

4 Count Sketches

In many cases where frequent items are sought the guaranty that $|f_i - g_i| \leq \varepsilon N = \varepsilon \sum_{j=1}^{n} f_j$ is insufficient. For example, if the item distribution is very skewed, a few most frequent items can correspond to most of the appearances in the stream. Thus, a more desirable guaranty would be of the form $|f_i - g_i| \leq \varepsilon N = \varepsilon \sum_{j=k+1}^{n} f_j$ for some prespecified k. Here we assume without loss of generality that the items are indexed in decreasing frequency order.

One idea from [6] suggests that this in possible by using $O(k\log(n/\delta))$ approximate counters. First, we create $3k$ different approximate counters and distribute elements between them using a hash function.

So, only with probability $1/3$ element a falls into the same sketch as one of the top k element. Therefore, with probability $2/3$, the sketch containing a will give a frequency approximation guaranty proportional to $\varepsilon \sum_{j=k+1}^{n} f_j$. Since this only happens with probability $2/3$ we must repeat the construction $O(\log(n/\delta))$ times and return the median of the results returned by the counters.

The second idea is to alter the approximate counters themselves by incorporating a random sign into the summation. That is, when element a is encountered, its counter is incremented by $s(a)$ where s is a hash function mapping items from the universe uniformly into $\{-1, 1\}$. This reduces the approximation error to be relative to $O(\sqrt{\sum_{j=k+1}^{n} f_j^2})$.

Algorithm 4 Count Sketch: Add

  Input: ε, A
  $C \leftarrow$ all zeros matrix of size $t \times b$
  for $i \in [N]$ **do**
    for $j \in [t]$ **do**
      $C[j, h_j(A_i)] = C[j, h_j(A_i)] + s_j(A_i)$
    end for
  end for
  Return: C

Algorithm 5 Count Sketch: Query

  Input: C, a
  Return: $\mathrm{median}_{j=1,\dots,t} \, C[j, h_j(a)]s(a)$

References

[1] Graham Cormode and S. Muthukrishnan. An improved data stream summary: the count-min sketch and its applications. *J. Algorithms*, 55(1):58–75, 2005.

[2] Gurmeet Singh Manku and Rajeev Motwani. Approximate frequency counts over data streams. In *VLDB*, pages 346–357, 2002.

[3] Jayadev Misra and David Gries. Finding repeated elements. Technical report, Ithaca, NY, USA, 1982.

[4] Erik D. Demaine, Alejandro López-Ortiz, and J. Ian Munro. Frequency estimation of internet packet streams with limited space. In *Proceedings of the 10th Annual European Symposium on Algorithms*, ESA '02, pages 348–360, London, UK, UK, 2002. Springer-Verlag.

[5] Richard M. Karp, Christos H. Papadimitriou, and Scott Shenker. A simple algorithm for finding frequent elements in streams and bags. *ACM Transactions on Database Systems*, 28:2003, 2003.

[6] Moses Charikar, Kevin Chen, and Martin Farach-colton. Finding frequent items in data streams. pages 693–703, 2002.