

Lecture 4: Home Assignment, Due Dec 3rd

Lecturer: Edo Liberty

Warning: This note may contain typos and other inaccuracies which are usually discussed during class. Please do not cite this note as a reliable source. If you find mistakes, please inform me.

1 Probabilistic inequalities

setup

In this question you will be asked to derive the three most used probabilistic inequalities for a specific random variable. Let x_1, \dots, x_n be independent $\{-1, 1\}$ valued random variables. Each x_i takes the value 1 with probability $1/2$ and -1 else. Let $X = \sum_{i=1}^n x_i$.

questions

1. Let the random variable Y be defined as $Y = |X|$. Prove that Markov's inequality holds for Y . Hint: note that Y takes integer values. Also, there is no need to compute $\Pr[Y = i]$.
2. Prove Chebyshev's inequality for the above random variable X . You can use the fact that Markov's inequality holds for any positive variable regardless of your success (or lack of it) in the previous question. Hint: $\text{Var}[X] = E[(X - E[X])^2]$.
3. Argue that

$$\Pr[X > a] = \Pr[\prod_{i=1}^n e^{\lambda x_i} > e^{\lambda a}] \leq \frac{E[\prod_{i=1}^n e^{\lambda x_i}]}{e^{\lambda a}}$$

for any $\lambda \in [0, 1]$. Explain each transition.

4. Argue that:

$$\frac{E[\prod_{i=1}^n e^{\lambda x_i}]}{e^{\lambda a}} = \frac{\prod_{i=1}^n E[e^{\lambda x_i}]}{e^{\lambda a}} = \frac{(E[e^{\lambda x_1}])^n}{e^{\lambda a}}$$

What properties of the random variables x_i did you use in each transition?

5. Conclude that $\Pr[X > a] \leq e^{-\frac{a^2}{2n}}$ by showing that:

$$\exists \lambda \in [0, 1] \text{ s.t. } \frac{(E[e^{\lambda x_1}])^n}{e^{\lambda a}} \leq e^{-\frac{a^2}{2n}}$$

Hint: For the hyperbolic cosine function we have $\cosh(x) = \frac{1}{2}(e^x + e^{-x}) \leq e^{x^2/2}$ for $x \in [0, 1]$.

answers

1.

$$\begin{aligned}
 E[Y] &= \sum_{i=0}^n \Pr[Y = i] \cdot i \\
 &= \sum_{i=0}^t \Pr[Y = i] \cdot i + \sum_{i=t+1}^n \Pr[Y = i] \cdot i \\
 &\geq \sum_{i=t+1}^n \Pr[Y = i] \cdot i \\
 &\geq \sum_{i=t+1}^n \Pr[Y = i] \cdot t \\
 &= t \cdot \Pr[Y > t]
 \end{aligned}$$

Therefore, $E[Y] \geq t \cdot \Pr[Y > t]$ which is Markov's inequality.

2. This is identical to the general proof of Chebyshev's inequality. We define $Z = (X - E[X])^2$. Since Z is positive we can use Markov's inequality for it and get:

$$\Pr[|X - E[X]| > t] = \Pr[Z > t^2] \leq \frac{E[Z]}{t^2} = \frac{\text{Var}[X]}{t^2}$$

Here we used that $E[Z] = E[(X - E[X])^2] = \text{Var}[X]$.

3. First transition:

$$\Pr[X > a] = \Pr[\lambda X > \lambda a] = \Pr[e^{\lambda X} > e^{\lambda a}] = \Pr[e^{\lambda \sum x_i} > e^{\lambda a}] = \Pr[\prod_{i=1}^n e^{\lambda x_i} > e^{\lambda a}]$$

These hold due to the monotonicity of multiplication by a positive constant and exponentiation. Now, using Markov's inequality on the last inequality we get:

$$\Pr[\prod_{i=1}^n e^{\lambda x_i} > e^{\lambda a}] \leq \frac{E[\prod_{i=1}^n e^{\lambda x_i}]}{e^{\lambda a}}$$

4. The first transition is true due to the independence of the variables x_i . This means that $e^{\lambda x_i}$ are independent. The second transition is due to all expectations of $e^{\lambda x_i}$ being equal which stems from x_i being identically distributed.

5. First, we compute the expectation of $e^{\lambda x_i}$

$$E[e^{\lambda x_i}] = \frac{1}{2}e^{\lambda} + \frac{1}{2}e^{-\lambda} = \cosh(\lambda) \leq e^{\lambda^2/2}$$

From the above we have that $\Pr[X > a] \leq e^{n\lambda^2/2 - \lambda a}$. Setting $\lambda = a/n$ we get $e^{n\lambda^2/2 - \lambda a} = e^{-\frac{a^2}{2n}}$ which concludes the proof.

2 Approximating the size of a graph

setup

In this question we will try to approximate the size of a graph. A graph $G(V, E)$ is a set of nodes $|V| = n$ and a set of edges $|E| = m$. Each edge $e \in V \times V$ is a set of two nodes which support it. We assume the graph is simple which means there are no duplicate edges and no self loops (i.e. an edge $e = (u, u)$). The degree of a node, $\deg(u)$, is the number of edges which it supports. More formally $\deg(u) = |\{e \in E | u \in e\}|$. The degree of each node in the graph is at least 1. The question refers to the following sampling procedure:

1. $e = (u, v) \leftarrow$ an edge uniformly at random from E .
2. with probability $1/2$
3. return u
4. else
5. return v

Throughout this question we assume that *i*) we can sample edges uniformly from the graph *ii*) that the number of edges m is known *iii*) that given a node u we can easily compute $\deg(u)$. The value of n , however, is unknown.

questions

1. Let $p(u)$ denote the probability that the sampling procedure returns a specific node, u . Compute $p(u)$ as a function of $\deg(u)$ and m . (Note: $\sum_{u \in V} \deg(u) = 2m$)
2. Let $f(u) = \frac{2m}{\deg(u)}$. Compute:

$$E_{x \sim smp}[f(x)]$$

where $x \sim smp$ denotes that x is chosen according to the distribution on the nodes generated by the above sampling procedure.

3. We say that a graph is d -degree-bounded if $\max_{u \in V} \deg(u) \leq d$. Show that for a d -degree-bounded graph:

$$\text{Var}_{x \sim smp}[f(x)] \leq dn^2$$

4. Let $Y = \frac{1}{s} \sum_{i=1}^s f(x_i)$ where x_i are nodes chosen independently from the graph according to the above sampling procedure. Compute $E[Y]$ **and** show that $\text{Var}[Y] \leq dn^2/s$.
5. Use Chebyshev's inequality to find a value for s such that for any d -degree-bounded graph and any two constants $\varepsilon \in [0, 1]$ and $\delta \in [0, 1]$:

$$\Pr[|Y - n| > \varepsilon n] < \delta.$$

s should be a function of d , ε and δ .

answers

1. A node is chosen only if an edge it is adjacent to is picked with probability $\frac{deg(u)}{2m}$ and then it is the node picked between the two. The first event happens with probability $\frac{deg(u)}{2m}$ since the edges are chosen uniformly at random. The second event happens with probability $1/2$ independently of the first event. This gives $p(u) = \frac{deg(u)}{2m} \cdot \frac{deg(u)}{2} = \frac{deg(u)^2}{4m}$.

2. By the definition to the expectation:

$$E_{x \sim smp}[f(x)] = \sum_{u \in V} p(u)f(u) = \sum_{u \in V} \frac{deg(u)}{2m} \cdot \frac{2m}{deg(u)} = \sum_{u \in V} 1 = n$$

3. We say that a graph is d -degree-bounded if $\max_{u \in V} deg(u) \leq d$. Show that for a d -degree-bounded graph:

$$\text{Var}_{x \sim smp}[f(x)] \leq E_{x \sim smp}[f^2(x)] = \sum_{u \in V} \frac{deg(u)}{2m} \left(\frac{2m}{deg(u)}\right)^2 = \sum_{u \in V} \frac{2m}{deg(u)}$$

Since $deg(u) \geq 1$ then $\sum_{u \in V} \frac{2m}{deg(u)} \leq \sum_{u \in V} \frac{2m}{1} = 2mn$. Also, since the graph is d -degree-bounded $2m = \sum_{u \in V} deg(u) \leq nd$ thus $2mn \leq dn^2$.

4. Y is the average of s independent copies of $f(x)$ and therefore, by linearity of the expectation, we have that $E[Y] = E[f] = n$. Moreover, Since the nodes x_i are chosen independently we have that $\text{Var}[Y] = \frac{1}{s^2} \sum_{i=1}^s \text{Var}[f(x_i)]$. Since $f(x_i)$ distribute identically and substituting $\text{Var}(x) \leq dn^2$ we get $\frac{1}{s^2} \sum_{i=1}^s \text{Var}[f(x_i)] \leq \frac{s}{s^2} dn^2 = dn^2/s$.

5. Since $E[Y] = n$ we get that the above holds if

$$\Pr[|Y - E[n]| > \varepsilon n] < \frac{\text{Var}[Y]}{\varepsilon^2 n^2} \leq \frac{dn^2/s}{\varepsilon^2 n^2} = \frac{d}{s\varepsilon^2}$$

The condition that $\frac{d}{s\varepsilon^2} \leq \delta$ holds for $s \geq \frac{d}{\delta\varepsilon^2}$

3 Approximate median

setup

Given a list A of n numbers a_1, \dots, a_n , we define the rank of an element $r(a_i)$ as the number of elements which are smaller than it. For example, the smallest number has rank zero and the largest has rank $n - 1$. Equal elements are ordered arbitrarily. The median of A is an element a such that $r(a) = n/2$ (rounded either up or down). An α -approximate-median is a number a such that:

$$n(1/2 - \alpha) \leq r(a) \leq n(1/2 + \alpha)$$

In this question we sample k elements uniformly at random *with replacement* from the list A . Let the samples be $\{x_1, \dots, x_k\} = X$. You will be asked to show that the median of X is an α -approximate-median of A .

questions

1. What is the probability the a randomly chosen element x is such that:

$$r(x) > n(1/2 + \alpha)$$

2. Let us define $X_{>\alpha}$ as the set of samples whose rank is greater than $n(1/2 + \alpha)$. More precisely, $X_{>\alpha} = \{x_i \in X | r(x_i) > n(1/2 + \alpha)\}$. Similarly we define $X_{<\alpha} = \{x_i \in X | r(x_i) < n(1/2 - \alpha)\}$. Prove that if $|X_{>\alpha}| < k/2$ and $|X_{<\alpha}| < k/2$ then the median of X is an α -approximate-median of A .
3. Let $Z = |X_{>\alpha}|$. Find t for which:

$$\Pr[Z \geq k/2] = \Pr[Z \geq (1 + t)E[Z]]$$

4. Bound from above the probability that $Z \geq k/2$ as tightly as possible. If you do so using a probabilistic inequality, justify your choice.
5. Compute the minimal value for k which will guarantee that $|X_{>\alpha}| < k/2$ **and** $|X_{<\alpha}| < k/2$ with probability at least $1 - \delta$.

answers

1. There are $n(1/2 - \alpha)$ elements for which $r(x) > n(1/2 + \alpha)$. Since the element is chosen uniformly, the probability of that happening is $(1/2 - \alpha)$.
2. First we note that the median of X cannot be either in $X_{>\alpha}$ or in $X_{<\alpha}$. This is simply because each of them includes less than half of the elements in X . Moreover, by the definitions of $X_{>\alpha}$ and $X_{<\alpha}$ we have:

$$n(1/2 - \alpha) \leq r(\text{median}(X)) \quad \text{and} \quad r(\text{median}(X)) \leq n(1/2 + \alpha)$$

which means that $\text{median}(X)$ is an α -approximate-median of A .

3. Since the probability of a sample being in $X_{>\alpha}$ is exactly $1/2 - \alpha$ and since we have k independent samples, $E[Z] = E[|X_{>\alpha}|] = k(1/2 - \alpha)$. Solving for t we get

$$(1 + t)E[Z] = k/2 \rightarrow (1 + t)(1/2 - \alpha) = 1/2 \rightarrow t = \frac{2\alpha}{1 - 2\alpha}$$

4. Since the value of Z is the sum of independent indicator variables we can apply Chernoff's inequality. Denoting $\mu = E[Z] = k(1/2 - \alpha)$ and $t = \frac{2\alpha}{1 - 2\alpha}$ we have:

$$\Pr[Z \geq k/2] = \Pr[Z \geq (1 + t)\mu] \leq e^{-\mu t^2/4}$$

5. Similarly to the the above we can argue that

$$\Pr[|X_{<\alpha}| \geq k/2] \leq e^{-\mu t^2/4}$$

From the union bound we have that the probability of the event that $|X_{<\alpha}| \geq k/2$ or that $|X_{>\alpha}| \geq k/2$ is at most the sum of their probabilities.

$$\Pr[|X_{<\alpha}| \geq k/2 \cup |X_{>\alpha}| \geq k/2] \leq \Pr[|X_{<\alpha}| \geq k/2] + \Pr[|X_{>\alpha}| \geq k/2] \leq 2e^{-\mu t^2/4}$$

Demanding that this failure probability is less than δ we guarantee success with probability at least $1 - \delta$. Substituting $\mu = k(1/2 - \alpha)$ and $t = \frac{2\alpha}{1 - 2\alpha}$ this is achieved for

$$2e^{-\mu t^2/4} < \delta \rightarrow k > \frac{4 \log(2/\delta)(1/2 - \alpha)}{\alpha^2}$$

4 Simple high capacity hashing

setup

In this question we try to evaluate the capacity of a special hash table. For simplicity, we assume that the hashed elements are a subset of $[N]$ ($[N]$ denotes the set $\{1, \dots, N\}$). The hash table consists of an array A of length n and L perfect hash functions $h_\ell : [N] \rightarrow [n]$. Throughout the exercise we assume the existence of perfect hash functions. That is, $\Pr[h(x) = i] = 1/n$ for all $x \in [N]$ and $i \in [n]$ independently of the values $h(x')$. For convenience we also assume that the entries in A are initialized to the value 0.

Algorithm 1 *Add(x)*

```
for  $\ell \in [L]$  do
  if  $A[h_\ell(x)] == 0$  or  $A[h_\ell(x)] == x$  then
     $A[h_\ell(x)] = x$ 
    Return Success
  end if
end for
Return Fail
```

Algorithm 2 *Query(x)*

```
for  $\ell \in [L]$  do
  if  $A[h_\ell(x)] == x$  then
    Return True
  else if  $A[h_\ell(x)] == 0$  then
    Return False
  end if
end for
Return False
```

questions

1. Argue the correctness of the hashing scheme. a) If an element was **successfully** added to the table by *Add(x)* it will be found by *Query(x)*. b) If an element was not added to the table by *Add(x)* it will not be found by *Query(x)*.
2. Assume that exactly m cells in the array are occupied. That is, m cells contain values $A[j] > 0$ and for the rest $A[j] = 0$. Given a new element x which is not stored in the hash table. What is the probability that location $h_1(x)$ in A is occupied.
3. What is the probability that procedure *Add(x)* fails for an element x not in the hash table? (here we still assume there are exactly m elements already in the table)
4. Assume we start with an empty hash table and insert m elements one after the other. Use the union bound to get a value for L for which *Add(x)* succeeds in **all** m element insertions with probability at least $1 - \delta$
5. Argue that the **expected** running time of both *Add(x)* and *Query(x)* is $O(1)$. That is, it does not depend on L .

answers

1. If $Add(x)$ returned “success” then for some ℓ we have $A[h_\ell(x)] = x$ and for any $\ell' < \ell$ it holds that $A[h_{\ell'}(x)] \notin \{0, x\}$. Therefore it will be found by $Query(x)$. Also, if x was not added then it cannot be found by $Query$ since it returns “True” only if $A[h_\ell(x)] = x$ for some ℓ .
2. Since x was not added and since h_1 is a perfect hash function then $\Pr[h_1(x) = i] = 1/n$ for all $i \in [n]$. Since there are m occupied cells this sums to $\Pr[A[h_1(x)] > 0] = m/n$.
3. Add fails only if for each to the $\ell \in [L]$ hash functions $A[h_\ell(x)] > 0$. Since they are chosen independently of each other we have

$$\Pr[Add(x) \text{ fails}] = (m/n)^L$$

4. Using the union bound we have that $\Pr[fail] \leq \sum_{i \in [m]} ((i-1)/n)^L$. This is because there are at most $i-1$ elements in the hash table when we insert the i th one. Computing this sum can be made simpler by bounding it with an integral.

$$\sum_{i \in [m]} ((i-1)/n)^L \leq \int_{t=1}^{m+1} ((t-1)/n)^L dt = \int_{t=0}^m (t/n)^L dt = \frac{1}{L+1} (m/n)^{L+1}$$

That said, even a bound as simple as $m(m/n)^L$ would have sufficed. For the sake of simplicity let us use the latter. We obtain that the failure probability is $m(m/n)^L \leq \delta$ if $L \geq \log(m/\delta)/\log(n/m)$. Note that the hash can contain millions of items and be at $\sim 80\%$ capacity and still $L \sim 100$.

5. Let us start with the expected running time of Add . Denote by $\ell^* = \min_\ell A[h_\ell(x)] = 0$. Clearly, the running is $O(\ell^*)$ since each lookup requires $O(1)$ time.

$$\mathbb{E}[\ell^*] = \sum_{\ell=1}^L \ell \Pr[\ell^* = \ell] \leq \sum_{\ell=1}^{\infty} \ell \left(\frac{m}{n}\right)^{\ell-1} \left(1 - \frac{m}{n}\right) = O(1)$$

This assumes the ratio between m and n is fixed. Regardless, this does not depend on L .

Now we argue the same about $Query$. If x has been added then $Query(x)$ takes the same amount of time that $Add(x)$ did at the time of insertion. If x has not been added then $Query$ returns *False* in the same amount of time it would have taken to run $Add(x)$. If both both cases it reduces the calculation above.